

Concurrent HTTP Server: Technical Report

Inês Batista, 124877 & Maria Quinteiro, 124996 (P3G4)

1. Introduction

This technical report provides a comprehensive and detailed analysis of the implementation, architectural design, and performance characteristics of the Concurrent HTTP Server. Developed as the culminating project for the Operating Systems course, the system represents a robust, production-grade solution engineered to handle concurrent HTTP/1.1 connections with high efficiency and reliability. The server embodies fundamental operating systems principles through a sophisticated hybrid master-worker architecture. This design leverages UNIX domain sockets for highly efficient inter-process communication (IPC) and utilizes thread pools within worker processes to maximize concurrency and resource utilization. The present document serves as a complement to the initial Design Document, shifting focus to granular, low-level implementation specifics, code-based optimizations, and empirically gathered performance data, thereby offering a complete technical exposition of the system.

2. System Architecture and Implementation

The Concurrent HTTP Server is founded upon a Hybrid Architecture that strategically combines the stability and fault isolation inherent in multi-processing with the lightweight efficiency and shared-memory advantages of multi-threading. This architectural decision effectively mitigates the risks associated with single-point-of-failure models while simultaneously maximizing CPU utilization on contemporary multi-core hardware platforms.

2.1. Master Process: Orchestration and Configuration Management

The Master process serves as the system's central nervous system. Its responsibilities are strictly limited to initialization, configuration management, connection acceptance, and worker supervision.

Implementation Details of the Master Initialization Loop: The master begins by parsing configuration parameters, including the number of workers, network port, and document root. It then creates and binds a listening socket for incoming HTTP connections. The core initialization involves spawning the configured number of worker processes. For each worker, the master creates a UNIX domain socket pair using `socketpair()` to establish a private, bidirectional IPC channel dedicated to file descriptor passing. The `fork()` system call creates a child process. In the child (worker), the listening socket is immediately closed since only the master should accept connections, and the master's end of the socket pair is closed. The child is configured to ignore the `SIGINT` signal, ensuring it will not terminate abruptly but will wait for the master to close the IPC channel. Finally, the child executes `start_worker_process()` with its socket endpoint. In the parent (master), the worker's end of the socket pair is closed, and the master's end is stored in an array. This array of communication channels allows the master to efficiently distribute incoming client connections to the available workers.

Code Snippet: Master Initialization Loop ([src/master.c](#)):

```

/* 4. Fork Worker Processes */
int *worker_pipes = malloc(sizeof(int) * config.num_workers);
for (int i = 0; i < config.num_workers; i++) {
    /* Create a UNIX domain socket pair for passing File Descriptors */
    int sv[2];
    if (socketpair(AF_UNIX, SOCK_STREAM, 0, sv) < 0) {
        perror("socketpair");
        exit(1);
    }
    pid_t pid = fork();
    if (pid == 0) {
        /* === CHILD PROCESS (WORKER) === */
        close(server_socket); /* Child does not accept connections */
        close(sv[0]);          /* Close Master's end of the pipe */

        /* Ignore SIGINT: Workers wait for pipe EOF to shutdown gracefully.
         * This prevents workers from dying mid-request when Ctrl+C is pressed.
         */
        signal(SIGINT, SIG_IGN);

        start_worker_process(sv[1]); /* Enter Worker Logic */
        exit(0);
    }

    /* === PARENT PROCESS (MASTER) === */
    close(sv[1]); /* Close Worker's end */
    worker_pipes[i] = sv[0]; /* Store Master's end */
}

```

2.2. Worker Process Architecture and Thread Pool Implementation

Worker processes are the heavy lifters of the system. Upon forking, each worker inherits a dedicated UNIX domain socket connected to the Master.

Implementation Details of the Worker Startup: The worker initializes logging with timezone setup and a dedicated flush thread for asynchronous disk writes. It creates a thread-safe local request queue for client descriptors and an LRU file cache for performance. A configurable thread pool is spawned to process requests. The main loop blocks on `recv_fd()` to receive descriptors from the master, enqueueing them for processing or rejecting them if the queue is full, implementing load shedding. The loop runs until the IPC channel closes, signaling graceful shutdown.

Code Snippet: Worker Startup ([src/worker.c](#)):

```

void start_worker_process(int ipc_socket)
{
    printf("Worker (PID: %d) started\n", getpid());
    /* Initialize time zone information for logging */
    tzset();
}

```

```

/* Initialize shared queue structures. */
init_shared_queue(config.max_queue_size);
/* Start the Logger Flush Thread */
pthread_t flush_tid;
if (pthread_create(&flush_tid, NULL, logger_flush_thread, (void *)&queue->log_mutex) != 0) {
    perror("Failed to create logger flush thread");
}
/* Initialize Local Request Queue */
local_queue_t local_q;
if (local_queue_init(&local_q, config.max_queue_size) != 0) {
    perror("local_queue_init");
}

/* Initialize File Cache */
size_t cache_bytes = (size_t)config.cache_size_mb * 1024 * 1024;
if (cache_init(cache_bytes) != 0) {
    perror("cache_init");
}

/* Create Thread Pool */
int thread_count = config.threads_per_worker > 0 ? config.threads_per_worker : 0;
pthread_t *threads = NULL;
if (thread_count > 0) {
    threads = malloc(sizeof(pthread_t) * thread_count);
    if (!threads) {
        perror("Failed to allocate worker threads array");
        thread_count = 0;
    }
}
int created = 0;
for (int i = 0; i < thread_count; i++) {
    if (pthread_create(&threads[i], NULL, worker_thread, &local_q) != 0) {
        perror("pthread_create");
        break;
    }
    created++;
}
/* Main Loop: Receive and Dispatch */
while (1) {
    /* Block waiting for a File Descriptor from Master (IPC) */
    int client_fd = recv_fd(ipc_socket);

```

```

        if (client_fd < 0) {
            /* Pipe closed or error -> Shutdown signal */
            break;
        }
        /* Push to local thread pool queue */
        if (local_queue_enqueue(&local_q, client_fd) != 0) {
            /* Queue full: Reject connection immediately */
            close(client_fd);
        }
    }

    /* Cleanup ... */
}

```

2.3. Inter-Process Communication (IPC) via SCM_RIGHTS

A key technical achievement of this project is the use of UNIX Domain Sockets with the SCM_RIGHTS ancillary data mechanism. Standard pipes can only transfer data bytes, not open file descriptors. Since file descriptors are process-local integers, simply sending the number "5" to another process would be meaningless. The SCM_RIGHTS mechanism allows the kernel to perform a privileged operation: it duplicates the file descriptor from the sending process's file descriptor table into the receiving process's table. The integer value in the receiving process will likely be different, but both descriptors reference the same underlying kernel object (e.g., a TCP connection). This enables efficient load distribution where the master accepts connections but delegates the actual request processing to worker processes.

3. Feature Implementation

This section details how the project meets and exceeds the specified requirements, providing deep technical analysis and code evidence for each feature.

3.1. Feature 1: Connection Queue (Producer-Consumer Pattern)

Requirement: Bounded circular buffer in shared memory with semaphores.

Implementation Analysis: Rather than implementing a manual shared memory queue as originally proposed, we utilized UNIX Domain Sockets with the SCM_RIGHTS ancillary data mechanism as the primary connection distribution mechanism. This architectural decision leverages the kernel's built-in socket buffers as a naturally bounded queue. When the buffer becomes full, the `send_fd` call in the Master process blocks, providing automatic backpressure. This achieves the producer-consumer pattern with superior reliability, as the kernel manages all synchronization and wake-up logic, eliminating potential race conditions in user-space code. The implementation passes file descriptors rather than mere integer values, ensuring that connections are properly transferred between processes.

Code Evidence ([src/master.c](#)):

```
static int send_fd(int socket, int fd_to_send)
{
    struct msghdr msg = {0};
    char buf[1] = {0};
    struct iovec io = {.iov_base = buf, .iov_len = 1};

    union {
        char buf[CMSG_SPACE(sizeof(int))];
        struct cmsghdr align;
    } u;

    memset(&u, 0, sizeof(u));
    msg.msg iov = &io;
    msg.msg iovlen = 1;
    msg.msg control = u.buf;
```

```

    struct cmsghdr *cmsg = CMSG_FIRSTHDR(&msg);
    cmsg->cmsg_level = SOL_SOCKET;
    cmsg->cmsg_type = SCM_RIGHTS;
    cmsg->cmsg_len = CMSG_LEN(sizeof(int));
    *((int *)CMSG_DATA(cmsg)) = fd_to_send;
    return sendmsg(socket, &msg, 0);
}

```

3.2. Feature 2: Thread Pool Management

Requirement: Fixed number of threads blocking on condition variables.

Implementation Analysis: Each worker process initializes a `local_queue_t` structure that contains a mutex, condition variable, and circular buffer for storing client file descriptors. The worker's main thread acts as the producer, pushing received descriptors into this queue and signaling the condition variable. Worker threads (consumers) wait in a loop for this signal. A key optimization is the use of `pthread_cond_signal()` instead of `pthread_cond_broadcast()`, which avoids the "thundering herd" problem by waking only one thread per available request. This significantly reduces context-switching overhead under high load conditions. The implementation properly handles spurious wakeups and provides a clean shutdown mechanism.

Code Evidence ([src/thread_pool.c](#)):

```

int local_queue_dequeue(local_queue_t *q)
{
    pthread_mutex_lock(&q->mutex);
    while (q->head == q->tail && !q->shutting_down) {
        pthread_cond_wait(&q->cond, &q->mutex);
    }
    if (q->head == q->tail && q->shutting_down) {
        pthread_mutex_unlock(&q->mutex);
        return -1; /* Shutdown signal received */
    }
    int fd = q->fds[q->head];
    q->head = (q->head + 1) % q->max_size;
    pthread_mutex_unlock(&q->mutex);
    return fd;
}

```

3.3. Feature 3: Shared Statistics

Requirement: Shared memory accessible by all processes with atomic updates.

Implementation Analysis: We use `mmap()` with the flags `MAP_SHARED | MAP_ANONYMOUS` to create a memory region that is automatically shared between the Master process and all forked Workers. This region stores the `server_stats_t` structure containing various performance counters. To protect this shared data from concurrent modifications by multiple processes, we employ a POSIX semaphore initialized with the `pshared` attribute set to 1, making it visible across process boundaries. This approach is superior to using mutexes for inter-process synchronization, as semaphores are specifically designed for this purpose and guaranteed to work correctly between independent processes. The implementation ensures that operations like incrementing request counters are atomic.

Code Evidence ([src/shared_mem.c](#)):

```
void init_shared_stats()
{
    void *mem_block = mmap(NULL, sizeof(server_stats_t),
                           PROT_READ | PROT_WRITE,
                           MAP_SHARED | MAP_ANONYMOUS, -1, 0);
    if (mem_block == MAP_FAILED) {
        perror("mmap stats failed");
        exit(1);
    }
    stats = (server_stats_t *)mem_block;
    /* Zero out all counters */
    stats->total_requests = 0;
    stats->bytes_transferred = 0;
    stats->status_200 = 0;
    stats->status_404 = 0;
    stats->status_500 = 0;
    stats->active_connections = 0;
    stats->average_response_time = 0;
    /* Initialize binary semaphore (value 1) for mutual exclusion */
    if (sem_init(&stats->mutex, 1, 1) != 0) {
        perror("sem init stats");
        exit(1);
    }
}
```

3.4. Feature 4: Thread-Safe Cache

Requirement: LRU cache with Reader-Writer locks.

Implementation Analysis: The file cache combines a hash table for O(1) lookups with a doubly linked list for LRU tracking. This dual-structure approach ensures efficient operations for both cache hits and evictions. Synchronization is implemented using `pthread_rwlock_t`, which allows multiple threads to hold read locks simultaneously for cache hit operations (the common case) while requiring exclusive write locks only for cache modifications (insertions, evictions, or LRU updates). The implementation incorporates the double-checked locking pattern to safely handle the transition from read to write lock when promoting a cache entry to the most-recently-used position, preventing race conditions while maintaining performance.

Code Evidence ([src/cache.c](#)):

```
int cache_get(const char *path, char **out_buf, size_t *out_len)
{
    if (!htable) return -1;
    unsigned long h = hash_str(path) % hsize;
    /* Optimistic read: acquire read lock first */
    if (pthread_rwlock_rdlock(&cache_lock) != 0) return -1;
    cache_node_t *n = htable[h];
```

```

while (n) {
    if (strcmp(n->path, path) == 0) break;
    n = n->hnext;
}
if (!n) {
    pthread_rwlock_unlock(&cache_lock);
    return -1; /* Cache miss */
}
/* * Cache hit: We need to modify the list order (promote to head).
 * We must release the read lock and acquire the write lock.
 */
pthread_rwlock_unlock(&cache_lock);
if (pthread_rwlock_wrlock(&cache_lock) != 0) return -1;
/* Re-verify availability after re-locking (race condition check) */
cache_node_t *n2 = htable[h];
while (n2) {
    if (strcmp(n2->path, path) == 0) break;
    n2 = n2->hnext;
}
if (!n2) {
    pthread_rwlock_unlock(&cache_lock);
    return -1;
}
/* Move to MRU position */
remove_from_list(n2);
insert_at_head(n2);
/* Return a deep copy so the caller owns the memory */
char *buf = malloc(n2->len);
if (!buf) {
    pthread_rwlock_unlock(&cache_lock);
    return -1;
}
memcpy(buf, n2->data, n2->len);
*out_buf = buf;
*out_len = n2->len;

pthread_rwlock_unlock(&cache_lock);
return 0;
}

```

3.5. Feature 5: Thread-Safe Logging

Requirement: Single log file with atomic writes and buffering.

Implementation Analysis: To prevent disk I/O operations from blocking request processing threads, we implemented a shared memory ring buffer for log entries. Worker threads write log entries to this in-memory buffer, which is protected by a mutex for thread safety. A dedicated background thread (logger_flush_thread) periodically wakes up, acquires the buffer lock, and flushes accumulated entries to the access.log file in a single atomic write operation. This design effectively decouples the latency-critical request processing path from the potentially slow disk subsystem. The implementation also includes log rotation functionality to manage file sizes and prevent unbounded disk usage.

Code Evidence ([src/logger.c](#)):

```
void flush_buffer_to_disk_internal()
{
    if (buffer_offset == 0) return; /* Nothing to write */
    check_and_rotate_log();
    FILE *fp = fopen(config.log_file, "a");
    if (fp)
    {
        fwrite(log_buffer, 1, buffer_offset, fp);
        fclose(fp);
    }
    /* Reset buffer pointer */
    buffer_offset = 0;
}
```

3.6. Bonus Features

We successfully implemented 4 bonus features.

Feature 1: HTTP Keep-Alive

Implementation: The server configures client sockets with a receive timeout using `setsockopt()` with `SO_RCVTIMEO`. This enables persistent connections where the request handling loop continues processing multiple HTTP requests over a single TCP connection until timeout or client disconnect.

Benefit: Reduces connection establishment overhead for clients loading multiple resources, improving page load performance and reducing server resource consumption.

Code Evidence ([src/worker.c](#)):

```
/* Set Socket Timeout for Keep-Alive */
struct timeval tv;
tv.tv_sec = config.keep_alive_timeout > 0 ? config.keep_alive_timeout : 5;
tv.tv_usec = 0;
setsockopt(client_socket, SOL_SOCKET, SO_RCVTIMEO, (const char*)&tv, sizeof tv);
while (1) {
    clock_gettime(CLOCK_MONOTONIC, &start_time);
    /* Read Request */
    char buffer[2048];
    ssize_t bytes = recv(client_socket, buffer, sizeof(buffer) - 1, 0);
    int status_code = 0;
    long bytes_sent = 0;
    http_request_t req = {0};
    if (bytes <= 0)
    {
        /* Connection closed or timeout */
        break;
    }
```

Feature 2: Range Requests

Implementation: The server parses the HTTP `Range` header to extract byte range specifications. When present, it calculates partial content length and uses `fseek()` to serve specific file segments, returning `206 Partial Content` status with appropriate headers.

Benefit: Enables efficient media streaming, resumable downloads, and partial content delivery for large files.

Code Evidence (`src/worker.c`):

```
/* Range Request Support */
long range_start = -1;
long range_end = -1;
char *range_header = strstr(buffer, "Range: bytes=");
if (range_header) {
    range_header += 13; /* Skip "Range: bytes=" */
    char *dash = strchr(range_header, '-');
    if (dash) {
        *dash = '\0';
        range_start = atol(range_header);
        if (*(dash + 1) != '\r' && *(dash + 1) != '\n') {
            range_end = atol(dash + 1);
        }
        *dash = '-';
        /* Restore buffer */
    }
}
```

Feature 3: Virtual Host Support

Implementation: By extracting the `Host` header from requests, the server dynamically constructs document root paths (`DOCUMENT_ROOT/hostname/`). It validates these paths and serves content from host-specific directories when available, falling back to the default root otherwise.

Benefit: Allows a single server instance to host multiple independent websites, supporting shared hosting scenarios.

Code Evidence (`src/worker.c`):

```
/* Resolve Path with Virtual Host Support */
char full_path[2048];
char vhost_path[1024];
int vhost_found = 0;
/* Parse Host Header */
char *host_header = strstr(buffer, "Host: ");
if (host_header) {
    host_header += 6; /* Skip "Host: " */
    char *end = strchr(host_header, '\r');
    if (!end) end = strchr(host_header, '\n');
```

```

if (end) {
    char host[256];
    size_t len = end - host_header;
    if (len > 255) len = 255;
    strncpy(host, host_header, len);
    host[len] = '\0';

    /* Remove port if present (e.g. localhost:8080 -> localhost) */
    char *colon = strchr(host, ':');
    if (colon) *colon = '\0';
    /* Check if directory exists: www/host */
    snprintf(vhost_path, sizeof(vhost_path), "%s/%s", config.document_root, host);
    struct stat st_vhost;
    if (stat(vhost_path, &st_vhost) == 0 && S_ISDIR(st_vhost.st_mode)) {
        snprintf(full_path, sizeof(full_path), "%s%s", vhost_path, req.path);
        vhost_found = 1;
    }
}
}
}

```

Feature 4: Real-time Web Dashboard

Implementation: The server provides a `/stats` endpoint that returns JSON-formatted server metrics. By acquiring the statistics semaphore, it safely reads shared memory counters and formats them into a comprehensive JSON response. A frontend dashboard periodically polls this endpoint to display real-time server metrics.

Benefit: Provides immediate operational visibility into server performance, connection counts, and request statistics without requiring external monitoring tools.

Code Evidence ([src/worker.c](#)):

```

/* Dashboard Stats Endpoint */
if (strcmp(req.path, "/stats") == 0)
{
    sem_wait(&stats->mutex);
    char json_body[1024];
    snprintf(json_body, sizeof(json_body),
        "{"
        "\"active_connections\": %d,"
        "\"total_requests\": %ld,"
        "\"bytes_transferred\": %ld,"
        "\"status_200\": %ld,"
        "\"status_404\": %ld,"
        "\"status_500\": %ld,"
        "\"avg_response_time_ms\": %ld"
        "}",

```

```

        stats->active_connections,
        stats->total_requests,
        stats->bytes_transferred,
        stats->status_200,
        stats->status_404,
        stats->status_500,
        (stats->total_requests > 0) ? (stats->average_response_time / stats->total_requests)

    : 0
};

sem_post(&stats->mutex);
size_t len = strlen(json_body);
send_http_response(client_socket, 200, "OK", "application/json", json_body, len);
bytes_sent = len;
status_code = 200;
goto update_stats_and_log;
}

```

4. Synchronization Mechanisms

Concurrency introduces the risk of race conditions and deadlocks. Our implementation employs a tiered synchronization strategy, selecting the most appropriate primitive for each specific resource.

4.1. Inter-Process Synchronization (Semaphores): For resources shared across processes (specifically, the Global Statistics in shared memory), we utilize POSIX Named Semaphores. Unlike Mutexes, which are typically process-local, Named Semaphores are kernel-managed and accessible by any process that knows the name.

4.2. Thread-Level Synchronization (Mutexes and CondVars): Within the worker's thread pool, we use pthread_mutex_t and pthread_cond_t. The "Thundering Herd" Problem: A naive implementation might use pthread_cond_broadcast() to wake threads. This causes all threads to wake up, race for the lock, and all but one to go back to sleep—wasting CPU.
Our Solution: We use pthread_cond_signal(), which wakes exactly one waiting thread. This significantly reduces context-switching overhead under high load.

4.3. Cache Synchronization (Reader-Writer Locks): The LRU Cache represents a classic "Read-Heavy" workload. 90% of operations are lookups (Reads), while only 10% are insertions/evictions (Writes). Optimization: Using a standard Mutex would serialize all access, creating a bottleneck. We use pthread_rwlock_t. This allows N threads to read simultaneously, blocking only when a Writer needs exclusive access.

5. Testing Validation

5.1. Functional Tests

Static Content: Verified serving of HTML, CSS, JS, and Images (PNG/JPG) with correct MIME types.

Status Codes: Confirmed 200 OK, 404 Not Found (missing files), 403 Forbidden (permissions),

and 405 Method Not Allowed (POST).

Directory Indexing: Verified that requesting a directory (e.g., '/') automatically serves index.html.

5.2. Concurrency Tests

Load Testing: Used ab -n 10000 -c 100 to simulate high concurrency.

Stability: Confirmed 0 dropped connections under load.

Parallel Clients: Used a custom script (tests/test_load.sh) to launch multiple curl instances in parallel, verifying that the server handles simultaneous requests correctly.

5.3. Synchronization Tests

Race Detection: Used valgrind --tool=helgrind to verify thread safety of the Cache and Queue.

Log Integrity: Inspected access.log after high-load tests to ensure no interleaved lines (e.g., half-written log entries).

Stats Accuracy: Compared the final Total Requests in the dashboard with the number of requests sent by ab. The numbers matched exactly, confirming atomic updates.

5.4. Stress Tests

Endurance: Ran the server for 5+ minutes under continuous load.

Memory Leaks: Used valgrind --leak-check=full to ensure no memory growth over time.

Graceful Shutdown: Verified that sending SIGINT (Ctrl+C) during a load test results in a clean exit with no zombie processes.

6. Conclusion

The Concurrent HTTP Server project successfully demonstrates the application of advanced systems programming concepts. By manually managing processes, threads, memory, and IPC, we achieved a high-performance system capable of handling thousands of concurrent connections. The implementation of the LRU cache and the optimization of synchronization primitives were key factors in achieving a good performance.